

Certified Position–Topology Trade-offs for Hierarchy-Constrained 1-D Layouts

Category: Research




Figure 1: ERA5 pressure lows over 59 days. (a) ST-MTM: unavoidable and avoidable order error after calibration. (b) Tree-preserving anchored map: certified space and hierarchy costs. (c) Relaxed map: actual position error and topological cost in hPa. The strips measure different quantities; lower panels show the pressure fields.

Abstract

Merge tree maps summarize time-varying scalar fields in a static image while preserving feature hierarchies. However, keeping each subtree contiguous can force features away from their positions along a spatial reference. We introduce Certificate-Driven Merge Tree Maps (CD-MTM), a method for diagnosing this conflict and controlling the position–topology trade-off. A positional certificate quantifies the least possible maximum reference error, separates hierarchy cost from packing cost, and identifies conflicting feature triples. We derive a fixed-order closed form and a display-resolution dynamic program for bounded-degree trees. A filling-independent topological bound and an optimal value-preserving filling characterize the exact attainable frontier for small trees. CD-MTM uses these results to retain or relax hierarchy and disclose both achieved position error and merge-level distortion. Across two ERA5 periods and an Australian wildfire sequence, hierarchy conflicts occur in 24–63% of frames under the selected extraction protocols. Relaxation reduces geometric motion-proxy reversals in several settings, with substantial, data-dependent topological costs. A 1024-row dendrogram example demonstrates the certificate’s wider applicability. The evaluation establishes computational guarantees and geometric trade-offs; human reading benefits remain to be tested, and full layout optimization remains limited by order enumeration.

Index Terms: Merge trees, hierarchy-constrained layout, time-varying scalar fields, static visualization, faithfulness, interleaving distance.

1 Introduction

Static visual summaries place an entire time sequence in one image, supporting comparisons across time without replaying an animation. Hovmöller diagrams summarize weather along a fixed direction [19]; MotionRugs summarize trajectories [5]; and merge tree maps linearize temporal scalar fields while preserving their merge trees [25, 9]. These designs place a spatial summary on one axis and time on the other. When a task calls for following features along a spatial direction, the question is whether their displayed positions support that reading.

Merge tree maps face a specific obstacle: their leaf order must keep every subtree contiguous so that the column encodes the input merge tree [25]. This hierarchy can conflict with geography. Two pressure lows may merge first because they share a trough, while a third low lies between them along the chosen reference direction (Fig. 2). A tree-preserving order must move at least one feature away from its reference. Similar constraints arise in clustered heatmaps [13] and geophylogenies [23], although hierarchy constraints are not required by all static summaries.

Projection quality measures describe the distortion of a produced view [44, 1]; such distortion can be difficult to recognize visually [29, 45]. ST-MTM measures distance stress and neighbourhood preservation [9]. These measures leave a design question unresolved: how much deviation is forced by the hierarchy, and how much can a different layout avoid? A second question follows: if position matters enough to relax the hierarchy, which merge relations change and by how much?

We address these questions with **Certificate-Driven Merge Tree Maps (CD-MTM)**. Given a task-defined reference, the method computes the least possible maximum positional deviation, identifies conflicts, and provides a controlled relaxation with an explicit topological cost. This follows the principle of defining visualization quality before optimizing it [30]. Our contributions are:

- A **positional certificate** for hierarchy-constrained interval layouts, including fixed-order and width-free forms, witness triples, and an exact display-resolution algorithm for bounded-degree trees (Sec. 4).
- A **position–topology trade-off**: a filling-independent lower bound, an optimal value-preserving filling, and an exact frontier for small trees (Sec. 5).
- **CD-MTM**, a diagnostic and relaxation pipeline with visual disclosure of both costs, evaluated on real and synthetic fields and a dendrogram example (Secs. 6 and 7).

The evidence supports a computational technique: it establishes what layouts can achieve and measures the resulting geometric trade-offs. The motion evaluation uses a geometric proxy, without claiming a demonstrated benefit to human readers. We first define the task and layout model, derive the certificates and trade-off, and then evaluate the pipeline against existing maps and simple reference-based alternatives.

2 Related Work

Static summaries of spatio-temporal data. Hovmöller and radius–time diagrams slice or aggregate one dimension; space-filling curves and pixel techniques linearize space [16, 50, 21]. MotionRugs, stable visual summaries, MoReVis, and GroupRugs summarize trajectories and regions [5, 48, 42, 40]. GroupRugs evaluates neighbourhood-based spatial quality under group contiguity but does not bound unavoidable conflict. Topology-based summaries include topological landscapes [46, 3], temporal merge tree maps (TMTM) [25], spatiotemporal merge tree maps (ST-MTM) [9], temporal treemaps and nested tracking graphs [24, 28]. TMTM favors overlap between columns; ST-MTM additionally fits geometric distances and widths. Both preserve the tree but do not certify positional deviation.

Hierarchy-constrained orders. Optimal leaf ordering [2] underlies clustered heatmaps [13]. Tree-constrained leaf orders against external positions appear in geophylogenies [23], in reordering a tree according to a given leaf order [6], and in displacement-optimized tanglegrams [20]; the objectives there are crossings, inversions, or summed displacement. Storyline layouts such as StoryFlow also order lines under hierarchy constraints and minimize crossings and wiggles [27]. We address a bottleneck objective (the maximum deviation), include widths, gaps, a canvas, and anchor eccentricity, and pair it with a topological cost.

Faithfulness, distortion, and trade-offs. Visualization mirages and misread embeddings motivate explicit distortion analysis [29, 45]. Trustworthiness, continuity, and distortion maps assess projections [44, 1, 38, 37]. Cartograms, data–spatial grid layouts, and spatially ordered treemaps trade or measure geometric fidelity [35, 43, 47]. Formal accounts of data–visual correspondence and task-specific projection distortion [22, 33, 34] assess produced views. We instead bound distortion over all admissible layouts and compute an attainable trade-off for small trees.

1-D placement. Without a hierarchy, minimizing the maximum displacement needed to separate intervals on a line is solvable in $O(n \log n)$ [26]; this is the free case of our certificate without canvas and eccentricity constraints. Separation-constrained quadratic programs are standard in constrained graph layout [11], and we use one to choose a layout within the certified budget.

Topological simplification. Persistence-based simplification removes low-persistence extrema and thereby simplifies merge trees [12, 7, 41]. It could also remove conflicts, but only by deleting features. Our flattening keeps every feature and removes a merge relation; its cost γ is the gap between two adjacent merge values, not the persistence of an extremum. Section 7.2 compares the two.

Merge tree comparison and visualization. Interleaving distances compare merge trees [31]; the ℓ_∞ -cophenetic metric of labelled trees is an interleaving distance [32]. Intrinsic and geometry-aware interleavings and Wasserstein metrics support comparison, tracking, and averaging [17, 49, 36]. ParkView encodes monotone interleavings between two ordered merge trees using structural guarantees [4]. Our problem concerns the distortion introduced when one input tree is rendered in a spatially referenced column, with both its positional and merge-level costs quantified. Persistent homology also evaluates and constructs projections [39, 10], and merge distortion compares cluster trees [14].

3 Problem and Motivation

3.1 Merge tree maps

For each time step t of a scalar field f_t on a 2-D domain, the merge tree T_t records how the components of sublevel sets (a *join tree*, e.g., for pressure minima) or superlevel sets (a *split tree*, e.g., for fire-intensity maxima) are born at extrema and merge at saddles. Its leaves are the features of interest. Leaf i has an extremum position $R_i \in \mathbb{R}^2$ and a support measure A_i (the area of its segment), and every pair of leaves has a merge value $f(\text{lca}(i, j))$. A merge tree map chooses, per time step, a leaf order, an interval per leaf, and an anchor inside the interval; fills the column by sampling values along leaf arcs and lowest-common-ancestor (LCA) paths [25, 9]; and stacks the columns over time. Because every subtree occupies a contiguous block, the merge levels read from the column equal those of T_t : the map preserves topology.

3.2 Tasks and design goals

We consider three candidate tasks: (T1) *compare features and merges*, using extremum values, support sizes, and merge levels; (T2) *locate features* along a specified spatial reference; and (T3) *compare approach or separation* over time. T1 motivates the topology-preserving design of merge tree maps. T2 and T3 motivate the geometric objectives of ST-MTM [9] and related 1-D projections [16]. These are analytical design targets; we have not established how readers perform them on CD-MTM.

They lead to three design goals. **G1: diagnose unavoidable displacement** before interpreting position (Sec. 4). **G2: retain features and quantify changed merge levels** when improving position (Sec. 5). **G3: disclose achieved error and topological cost** beside the map so that its operating point is visible (Sec. 6). The evaluation tests G1 and G2 computationally; the case study illustrates G3 without serving as a user study.

Reference coordinate. The task supplies a scalar reference $q_i = \phi(R_i)$ for feature i . All certificates hold for any given reference. For the retrospective temporal summaries evaluated here, we use the principal direction of tracked feature motion and test fixed directions for robustness. This reference measures location and motion along one direction; it cannot preserve all 2-D distances. TMTM and ST-MTM optimize different goals. Our certificate quantifies the deviation that their tree-order constraint necessarily incurs for the specified reference.

Why the maximum deviation. We bound the maximum rather than an average deviation. A reader does not know which feature is misplaced, so a guarantee must hold for every feature. Averages and stress hide the few misplaced features on which motion readings reverse (Sec. 7.3). And the maximum admits a closed form, an exact algorithm, and a clean interaction with topology (Sec. 5). Average-

type measures remain useful for choosing among layouts within the certified budget, which is how we use them (Sec. 6).

Why optimize the bound? A reference-anchored layout minimizes maximum deviation under the hierarchy, then balances secondary quality criteria within a small slack. On ERA5 it still exceeds 2% of the axis in 34% of frames and reverses 19.2% of clear geometric motion-proxy readings, compared with 21.2% for ST-MTM. This motivates separating unavoidable hierarchy cost from layout choices, rather than attributing all error to an optimizer.

3.3 Model

Let the hierarchy be a rooted tree T with leaves $1, \dots, n$. Leaf i has a reference q_i , a width $w_i = cA_i$ with a scale c that is fixed over time, and an eccentricity radius $h_i = \rho w_i/2$ with $\rho \in [0, 1]$. Let $g > 0$ be a minimum gap and $[a, a+L]$ the canvas. A *layout* (π, z, u) consists of a leaf order π , interval centres z , and anchors u that satisfy

- (H) $\pi \in \Pi(T)$: the leaf set S_v of every node v is contiguous in π ;
- (D) $z_y - z_x \geq (w_x + w_y)/2 + g$ for π -consecutive leaves x, y ;
- (C) $a + w_i/2 \leq z_i \leq a + L - w_i/2$;
- (E) $|u_i - z_i| \leq h_i$.

Legal orders arise by permuting the blocks of children at every node, so $|\Pi(T)| = \prod_v \deg(v)!$. We distinguish three fidelities: *position*, the deviation $u_i - q_i$; *size*, widths proportional to support measure and comparable over time; and *topology*, the merge levels shown by the rendered map (Sec. 5). TMTM preserves topology and absolute size but does not control position; ST-MTM preserves topology, reflects relative distances, and fits target width shares; the reference-anchored layout targets position under (H). Our analysis concerns position and topology at fixed size.

4 A Positional Certificate

4.1 Closed form and decomposition

Definition 1. For an order π , $\tau(\pi) = \min_{z,u} \max_i |u_i - q_i|$ subject to (D), (C), (E). The certificate is $\tau^*(T) = \min_{\pi \in \Pi(T)} \tau(\pi)$, and τ_{free} is the same minimum over all $n!$ permutations. The hierarchy cost is $H = \tau^* - \tau_{\text{free}}$ and the space cost is τ_{free} .

τ^* is a property of the data and the hierarchy, not of an algorithm: no layout that keeps every subtree contiguous can deviate by less, and some layout attains it. For a fixed order it can be written down directly.

Proposition 2 (Closed form). *Number the leaves along π and let $s_k = (w_k + w_{k+1})/2 + g$, $S_{ij} = \sum_{k=i}^{j-1} s_k$, $D_k = \sum_{m < k} w_m + w_k/2 + (k-1)g$, and $E_k = \sum_{m > k} w_m + w_k/2 + (n-k)g$. If $\sum_i w_i + (n-1)g \leq L$, then*

$$\tau(\pi) = \max \left\{ 0, \max_{i < j} \frac{1}{2} (q_i - q_j + S_{ij} - h_i - h_j), \max_k (a + D_k - q_k - h_k), \max_k (q_k - h_k - a - L + E_k) \right\}.$$

Eliminating the anchors turns each constraint into a box window $|z_k - q_k| \leq \tau + h_k$ for the centre. A chain with lower-bounded separations and box windows is feasible if and only if its earliest placement stays inside every window, and expanding the resulting inequalities gives the formula. The pair term says that two leaves in inverted reference order must share the overlap of their windows; the other two terms account for the canvas ends. The formula agrees with a linear-programming solver to $2 \cdot 10^{-14}$ on our data.

Algorithm 1 Certificate τ^* at display resolution (Theorem 4).

Require: hierarchy T ; pixel widths W_i , gap G , references q_i , radii h_i ; N pixels; tolerance ϵ

- 1: **function** Earliest(S, τ) ▷ array over starts $x = 0, \dots, N$
- 2: **if** S is a leaf i **then**
- 3: $lo_i \leftarrow \max(0, \lceil q_i - \tau - h_i - W_i/2 \rceil)$, $hi_i \leftarrow \min(N - W_i, \lfloor q_i + \tau + h_i - W_i/2 \rfloor)$
- 4: **return** $x \mapsto \max(x, lo_i) + W_i + G$ if $\max(x, lo_i) \leq hi_i$,
 else ∞
- 5: $E_c \leftarrow \text{Earliest}(c, \tau)$ for every child c of S
- 6: $F_0 \leftarrow \text{id}$; $F_M \leftarrow \min_{c \in M} E_c \circ F_{M \setminus \{c\}}$ for all subsets M of children
- 7: **return** $F_{\text{all children}}$ ▷ binary: $\min(E_B \circ E_A, E_A \circ E_B)$
- 8: $\tau^* \leftarrow$ smallest τ with Earliest(root, τ)(0) $< \infty$, by bisection to ϵ

Proposition 3 (Monotonicity, decomposition, witnesses). *Flattening a node, i.e., attaching its children to its parent, enlarges Π and never increases τ^* ; flattening all nodes yields τ_{free} . Hence $H \geq 0$ is exactly the price of keeping every subtree contiguous. If $i, j \in S_v$, $k \notin S_v$, and $q_i < q_k < q_j$, then $\tau^* \geq \frac{1}{2} \min(q_k - q_i, q_j - q_k)$.*

The last statement holds because contiguity of S_v forces u_k outside $[u_i, u_j]$, so one of the three anchors must cross a reference that lies between the others. Such a *witness triple* names the conflicting features. In Fig. 2, $(A, B; C)$ gives the bound 1.4, and the widths raise τ^* to 1.85. On the real datasets, the best witness explains a median of 92–94% of τ^* in conflict frames, so a conflict is usually caused by one identifiable triple of features.

4.2 Exact computation at display resolution

Enumerating $\Pi(T)$ is exponential in the number of internal nodes. We therefore compute τ^* at the resolution of the display. Discretize the canvas into N pixels with integer widths W_i and gap G , and let $E_S(x)$ be the earliest start that is available after placing all leaves of subtree S in some legal order, none starting before x .

Theorem 4 (Exact dynamic program). *For a leaf, $E_i(x) = \max(x, lo_i) + W_i + G$, or ∞ if $\max(x, lo_i) > hi_i$, where $[lo_i, hi_i]$ is the window of Proposition 2. For a binary node, $E_S = \min(E_B \circ E_A, E_A \circ E_B)$; for an m -ary node, a subset recursion over its children. A deviation τ is feasible iff $E_{\text{root}}(0) < \infty$. Bisection yields τ^* to a tolerance ϵ in $O(nN \log(N/\epsilon))$ time for binary hierarchies—pseudo-polynomial in N —and $O(2^m mN)$ per m -ary node.*

Correctness rests on an earliest-completion argument. For a fixed order, the greedy earliest placement minimizes every start; hence a block influences the rest of the chain only through its earliest next start, and a smaller next start is never worse. The functions E_S are monotone arrays over $N+1$ start positions, and composing them is an array lookup. Algorithm 1 lists the computation.

Proposition 5 (Certified lower bound). *If widths and gap are integer multiples of the pixel size λ and the canvas is grid-aligned, then $0 \leq \tau_{\text{disc}}^* - \tau_{\text{cont}}^* \leq \lambda/2$. For real widths, rounding widths and gap down while keeping the original slacks gives $\tau_{\text{cont}}^* \geq \tau_{\text{disc, relax}}^* - \lambda/2$.*

Every continuous layout remains feasible for the relaxed discrete problem after shifting its starts to the pixel grid, which costs at most half a pixel; reported values also subtract the bisection tolerance. All values of the lower frontier Φ reported in Sec. 7 use this certified bound. On 300 random hierarchies the program matches brute-force enumeration within 0.5 pixels; on all ERA5 and Ring frames, whose widths are not integer pixels, within 0.6 pixels at about 2 ms per frame; and a 4096-leaf hierarchy on 16384 pixels takes 3.6 s. The scalable part is the certificate; choosing a layout within the budget still enumerates candidate orders (Sec. 6).




Figure 2: Why hierarchy and space conflict. (a) A scalar field with three minima; A and B share a trough, and C lies between them along the reference direction (dotted: projections q). (b) Its join tree $((A,B),C)$, computed from the field; γ_v is the gap between the A-B merge and the root. (c) Sorting by the reference splits the subtree $\{A,B\}$; the best layout that keeps it contiguous must deviate by τ^* , shared by B and C (arrows). (d) The resulting column of a merge tree map: leaf values at the anchors and the merge values between them, so that the column shows the merge tree.

4.3 A width-free certificate

Other merge tree maps use other width models: TMTM allocates one pixel per sample, and ST-MTM fits target widths. To bound them in their own models, we drop widths, gaps, eccentricity, and canvas, which removes every constraint except the order.

Proposition 6 (Width-free certificate). *Let $\tau_{\text{pt}}(\pi) = \max\{0, \max_i \text{before } j(q_i - q_j)/2\}$ and $\tau_{\text{pt}}^* = \min_{\pi \in \Pi(T)} \tau_{\text{pt}}(\pi) \leq \tau^*$. For every 1-D map whose anchor order is hierarchy-consistent, with any width model, and for every non-decreasing calibration c of its axis onto the reference axis, $\max_i |c(u_i) - q_i| \geq \tau_{\text{pt}}^*$.*

Two anchors in inverted order cannot both be closer than half their reference gap, whatever the calibration, and $\tau_{\text{pt}}(\pi)$ is attained by $c(u_k) = \frac{1}{2}(\max_{i \leq k} q_i + \min_{j \geq k} q_j)$. The width-free certificate therefore separates, for any existing map, the deviation that its hierarchy forces from the deviation that its layout chose (Fig. 1a).

5 The Price of Topology

5.1 Measuring topology on a 1-D map

Relaxing (H) can reduce the positional error, but the map then no longer shows the merge tree exactly. To quantify the change, consider any function ψ on the line (a *filling*) with leaf anchors u_i . In a join tree, the 1-D merge level of leaves i, j is $m_\psi(i, j) = \max_{[u_i, u_j]} \psi$, because the components of sublevel sets on a line are intervals (split trees are symmetric). We define the *topological distortion*, the largest error of an anchor-to-anchor merge level,

$$d_{\text{top}} = \max_{i \neq j} |m_\psi(i, j) - f(\text{lca}_T(i, j))|.$$

When the local minima of the column are exactly the labelled anchors, with matching leaf values, d_{top} is the ℓ_∞ -cophenetic distance between T and the merge tree of the column, i.e., their labelled interleaving distance [32]. It is zero exactly when the column shows the merge tree. Throughout, critical values are distinct (so $\gamma_v > 0$), statements are given for join trees, and split trees follow by negating f .

5.2 A lower bound for every filling

Theorem 7 (Topological price of non-contiguity). *If the leaf set of a non-root node v is not contiguous in a 1-D map, then for every filling ψ , $d_{\text{top}} \geq \gamma_v/2$ with $\gamma_v = f(p(v)) - f(v)$, where $p(v)$ is the parent of v . The factor $1/2$ is tight.*

Proof. Choose $i, j \in S_v$ and $k \notin S_v$ with u_k between u_i and u_j . Interval inclusion gives $m_\psi(i, k) \leq m_\psi(i, j)$, whereas $m_T(i, j) \leq f(v)$ and $m_T(i, k) \geq f(p(v))$. Hence $\gamma_v \leq m_T(i, k) - m_T(i, j) = [m_T(i, k) - m_\psi(i, k)] + [m_\psi(i, k) - m_\psi(i, j)] + [m_\psi(i, j) - m_T(i, j)] \leq 2d_{\text{top}}$. For tightness, take $T = ((i, j)_v, k)$ with $f(k) \leq f(v) + \gamma_v/2$, anchors in order i, k, j , and both barriers at $f(v) + \gamma_v/2$. $\square$

Figure 3 illustrates the theorem and the optimal filling below on the example of Fig. 2. The theorem constrains the map, not the algorithm: it holds for TMTM, ST-MTM, and any future method that breaks contiguity. Weak merges (small γ_v) are cheap to break; strong merges are expensive.

5.3 The optimal filling

Theorem 7 bounds every filling from below. For a fixed order, the best *value-preserving* filling, i.e., one that shows every extremum at its true value, can be given in closed form. Number the leaves along the map, keep the true leaf values at the anchors, and let $P_{ij} = f(\text{lca}(i, j))$ for $i < j$, $C_k = \min_{i \leq k < j} P_{ij}$, and $\ell_k = \max(f(k), f(k+1))$.

Proposition 8 (Optimal filling). *Over value-preserving fillings, $\inf_\psi d_{\text{top}} = \delta^*(\pi) = \max\{0, \max_k (\ell_k - C_k), \max_{i < j} \frac{1}{2}(P_{ij} - \max_{i \leq k < j} C_k)\}$. It is attained by the barrier filling, whose maximum between anchors k and $k+1$ is $b_k = C_k + \delta^*(\pi)$; if every anchor must remain a strict minimum, it is approached arbitrarily closely by raising barriers with $b_k = \ell_k$ by any $\epsilon > 0$. Moreover, $\delta^*(\pi) = 0$ iff $\pi \in \Pi(T)$.*

Any filling acts on merge levels only through its barrier heights b_k between consecutive anchors, with $m_\psi(i, j) = \max_{i \leq k < j} b_k$. The upper constraints $b_k \leq P_{ij} + \delta$ collapse to $b_k \leq C_k + \delta$, and raising a barrier only helps the lower constraints, so $b_k = C_k + \delta$ is optimal; feasibility then reduces to the two inequalities in the formula. The term $\ell_k - C_k$ reveals a second mechanism next to Theorem 7: a shallow extremum placed between two leaves that merge deeply raises their 1-D merge level at least to its own value, whatever value-preserving filling is used. If a filling may also move extremum values, which the labelled interleaving distance permits, this term halves to $\frac{1}{2}(\ell_k - C_k)$; we keep extrema at their true values because readers read depths from them. The pair term has the structure of ℓ_∞ fitting by ultrametrics [15, 8], here constrained to the ultrametrics that a line with fixed leaf values can realize. The LCA-path filling of existing maps is not optimal; on our data it pays up to twice

δ^* . The barrier filling changes a rendered column only between anchors, by clipping each segment to b_k and raising its existing peak to b_k if necessary, so no new extrema appear; our renderer raises barriers by about 10^{-6} relative to the field values above adjacent leaf values, so every anchor stays a strict extremum and every feature remains visible.

5.4 The position–topology frontier

Corollary 9 (Certified lower frontier). *Let $\Phi_T(\delta)$ be τ^* of T with every node of $\gamma_v \leq 2\delta$ flattened. Every layout satisfying (D), (C), (E), with any filling, obeys $\max_i |u_i - q_i| \geq \Phi_T(d_{\text{top}})$.*

By Theorem 7, every non-contiguous node of a layout with distortion d_{top} has $\gamma_v \leq 2d_{\text{top}}$, so the layout is legal for the flattened tree, whose certificate is $\Phi_T(d_{\text{top}})$ (Proposition 3). No 1-D map can simultaneously keep the positional error below $\Phi_T(\delta)$ and the topological distortion at or below δ . Φ_T is a step function with breakpoints at $\gamma_v/2$, and we evaluate it per frame with the certified lower bound of Proposition 5. It is generally not attainable. Computing it with the DP remains exponential in the largest child degree after flattening, which can be much larger than in the input tree. For small trees, the attainable frontier itself can be computed exactly.

Theorem 10 (Exact position–topology frontier). *A value-preserving map with order π has $\max_i |u_i - q_i| \geq \tau(\pi)$ and $d_{\text{top}} \geq \delta^*(\pi)$, and both bounds are attained, or approached arbitrarily closely, simultaneously. Hence the achievable pairs over all value-preserving 1-D maps have the lower envelope $F_T(\delta) = \min\{\tau(\pi) : \pi \in S_n, \delta^*(\pi) \leq \delta\}$, with $F_T(0) = \tau^*$, $F_T(\infty) = \tau_{\text{free}}$, and $\Phi_T \leq F_T$.*

Positions enter only through $\tau(\pi)$ (Proposition 2) and merge levels only through $\delta^*(\pi)$ (Proposition 8). The barrier filling prescribes values only between anchors, which remain separated by the gap g , so the position-optimal layout and the optimal filling can be combined. We evaluate F_T by enumerating all $n!$ orders with both closed forms, which covers every frame in our data ($n \leq 11$).

6 CD-MTM: From Certificates to Maps

CD-MTM turns the preceding guarantees into a four-stage pipeline: establish a reference, diagnose hierarchy cost, select a layout with optional relaxation, and render both the map and its cost disclosures (Algorithm 2). The inputs are merge trees, feature supports, and correspondences; the output is a static temporal map with per-frame positional and topological measurements. The tree-preserving variant retains all merges. The relaxed variant retains all features and changes selected merge relations.

6.1 Reference and conflict diagnosis

Reference. For retrospective summaries, d is the unit eigenvector of the largest eigenvalue of $\sum \Delta R_i \Delta R_i^T$, summed over matched extremum displacements. We use $q_i = \langle R_i, d \rangle$ on the fixed projected domain extent. The same rule is used for every dataset and is invariant to rotations of the input coordinate system. It captures 0.71 of the squared tracked motion in ERA5 2013/14 and 0.89 in 1999/2000. It uses the complete sequence, so the method as evaluated is an offline summary; prospective use would require a direction specified or fitted on earlier data.

Gate. For the small trees in our full-layout experiments, we compute the continuous τ^* and τ_{free} by order enumeration and relax only if $H > \theta = 2\%$ of the axis. This is an operating threshold, not a perceptually validated limit.

6.2 Relaxation and layout optimization

Threshold and prune. In a frame with $H > \theta$ we flatten all nodes with $\gamma_v \leq \kappa$, then undo flattenings, strongest first, while keeping a rounded display-model score $\hat{\tau}$ within half a pixel of its value with

Dataset	steps	Δt	leaves	tracks	tree / range
ERA5 1999/2000	118	12 h	1–8	135	join / 95.1 hPa
ERA5 2013/14	124	12 h	1–7	147	join / 112.1 hPa
Wildfire 2019	30	2 d	6–11	76	split / 4.2
Ring	40	–	1–7	19	split / 12.7
Gaussians	64	–	3–4	4	split / 1.5

Table 1: Datasets. Leaves: merge-tree leaves per time step; range: value range of the field (fire: $\log(1 + \text{FRP})$).

all admissible nodes flattened. The comparison uses this fixed initial value, so tolerance cannot accumulate across removals. Widths and gap are rounded to a 4096-pixel grid for this heuristic score; the final positional budget uses the exact continuous optimum. Joint flattening handles plateaus on which a greedy one-node policy stalls (Sec. 7.4). The threshold κ selects nodes by merge gap; it does not bound the resulting topological distortion, which Proposition 8 computes. The policy has no approximation guarantee.

Layout. Within the budget $B_t = \tau^*(T') + \beta$ of the possibly relaxed tree T' , we choose the order and positions that minimize a weighted sum of distance stress, reference error, eccentricity, and a motion residual that asks each tracked anchor to move like its reference since the previous frame, subject to (D), (C), (E) and $|u_i - q_i| \leq B_t$. This is one quadratic program per candidate order; candidate orders are those with $\tau(\pi) \leq B_t$. The weights of reference error, distance stress, motion residual, and eccentricity are fixed at 4, 1, 0.5, and 0.1, each normalized by its number of terms.

6.3 Rendering and visual disclosure

Columns are filled along leaf arcs and LCA paths. Relaxed columns receive the barrier correction, attaining $\delta^*(\pi)$ up to the strict-extremum clearance. Three disclosures address G3. A *certificate strip* separates space and hierarchy cost and marks θ (Fig. 1b). A *witness overlay* identifies two leaves of a subtree with a bar and the intervening reference feature with a ring (Fig. 6a). An *achieved-error and topology strip* reports actual positional error and merge-level distortion in data units (Fig. 1c). The certificate is a lower bound on unavoidable error, whereas achieved error describes the selected layout; displaying them separately prevents treating a small certificate as a guarantee about an arbitrary map. Strips preserve colour for field values and position for the reference. Their readability and utility require perceptual validation.

Constants. All constants are relative to the canvas: the width scale maps the domain area to half the axis, $g = L/240$, $\beta = L/120$, $\rho = 0.5$, and $\theta = 2\%$ of the axis. On a 400-pixel map, θ is 8 pixels; we use it as an operational tolerance without claiming a perceptual detection threshold. The same values are used for all datasets; Sec. 7.2 reports $\theta = 1\%$ and 5% . The merge-gap threshold κ is not a constant but the user’s choice of operating point on the position–topology trade-off; it bounds the gaps of flattened merges, not the resulting distortion δ^* , which the strips report; we report $\kappa \in \{2\%, 20\%, \infty\}$ of the value range and use 20% in figures.

7 Evaluation

The evaluation follows the contributions. **RQ1** tests whether the certificate identifies consequential hierarchy conflicts in the selected data. **RQ2** measures geometric motion-proxy error, separating projected and 2-D references. **RQ3** measures the topological cost of relaxation and compares its operating points with the exact frontier. **RQ4** tests certificate transfer to a hierarchical pixel display. These computational experiments and illustrative cases do not evaluate human task performance.

7.1 Experimental protocol

Pressure data (Tab. 1). We use ERA5 mean sea-level pressure over the North Atlantic and Europe (25–80°N, 40°W–50°E) [18], for




Figure 3: The price of topology on the example of Fig. 2 (columns as in Fig. 2d; triangles: references). (a) Keeping the merge tree reads every merge correctly but misplaces B and C by τ^* . (b) Placing every anchor at its reference separates A and B; with the LCA-path filling the A–B merge is read at the root, an error of γ_v . (c) The barrier filling of the same order splits the error: $d_{top} = \gamma_v/2$, the lower bound of Theorem 7. (d) The resulting trade-off: no 1-D map lies in the grey region below the certified bound Φ ; (a) and (c) lie on it, (b) does not.

Algorithm 2 Certificate-driven merge tree map. Constants: $\theta = 2\%$ of the axis; canvas-relative g, β, ρ ; budget κ .

Require: merge trees T_i with leaf extrema R_i and supports A_i ; feature tracks

- 1: $d \leftarrow$ principal direction of tracked extremum displacements;
- 2: $q_i \leftarrow \langle R_i, d \rangle$ scaled to the axis
- 3: **for** $t = 1, \dots, T$ **do**
- 4: $\tau^* \leftarrow \min_{\pi \in \Pi(T_i)} \tau(\pi)$; $\tau_{free} \leftarrow \min_{\pi \in S_n} \tau(\pi)$; $H_t \leftarrow \tau^* - \tau_{free}$
- 5: $T' \leftarrow T_t$
- 6: **if** $H_t > \theta$ **then** $\triangleright$ threshold and prune
- 7: $W \leftarrow \{v : \gamma_v \leq \kappa\}$; $b_0 \leftarrow \widehat{\tau}(\text{flat}(T_t, W))$
- 8: **for** $v \in W$ by decreasing γ_v **do**
- 9: **if** $\widehat{\tau}(\text{flat}(T_t, W \setminus \{v\})) \leq b_0 + \text{half a pixel}$ **then**
- 10: $W \leftarrow W \setminus \{v\}$
- 11: $T' \leftarrow \text{flat}(T_t, W)$
- 12: $B_t \leftarrow \tau^*(T') + \beta$ $\triangleright$ error budget
- 13: $(\pi, z, u) \leftarrow$ over $\pi \in \Pi(T')$ with $\tau(\pi) \leq B_t$, minimize distance stress, reference error, eccentricity, and motion residual w.r.t. $t-1$ subject to (D)(C)(E) and $|u_i - q_i| \leq B_t$ $\triangleright$ one QP per order
- 14: fill column t along leaf arcs and LCA paths; clip barriers to $C_k + \delta^*(\pi)$ (Prop. 8)
- 15: strips: $\tau^*, H_t, \max_i |u_i - q_i|, \delta^*(\pi)$

1999-11-17 to 2000-01-14 and December 2013 to January 2014. Fields are sampled every 12 hours on an equal-area 49×49 grid with 250 km smoothing; features are join-tree minima. A narrower window creates boundary minima from flanks of lows outside it: 62% of its leaves lie on the boundary. We therefore smooth on a grid enlarged by five cells per side (three to the north), retain only minima of the enlarged field that lie inside the window, and raise other basins to their spill levels. Retained leaves are minima of the processed enlarged field; they are not independently validated cyclone centres. Only 6–7% lie within one cell of the boundary. The buffered window fits within the downloaded 20–90°N, 50°W–60°E domain.

Fire and synthetic data. Australian fire radiative power at 10 km resolution [16], November–December 2019, is aggregated to two-day frames on a 64×64 grid of $\log(1 + \text{FRP})$. A weak background removes zero plateaus. We use the smallest smoothing yielding at most 12 maxima per frame so that exact-frontier enumeration is feasible; split-tree maxima represent processed fire hot

spots. Ring [25, 16] and a replica of ST-MTM’s three-Gaussian scene are synthetic controls. Extraction sensitivity is reported in RQ1; these protocols define the analyzed features and limit generalization of prevalence rates.

Methods and comparison roles. TMTM and our reimplementations of ST-MTM [9] represent existing tree-preserving maps. ST-MTM uses its published Ring parameters and adapted defaults elsewhere; no original implementation was available for comparison. CD-MTM is evaluated with the hierarchy kept (A) and relaxed (R) at $\kappa \in \{2\%, 20\%, \infty\}$ of the value range. A *sorted, packed* baseline uses reference order with R’s widths, gaps, canvas, and secondary objective; its optimal filling measures the cost of dropping hierarchy entirely. An *exact-projection* baseline places every anchor at its reference, removing packing and hierarchy constraints to expose projection loss. All methods share extracted trees, supports, and correspondences. A fixed-direction version of A is a robustness control, rather than a separate proposed method.

Checks. Every layout is checked for legality, widths, gaps, canvas, eccentricity, and budget, and every unrelaxed frame for exact equality of the rendered merge tree. Theorem 7 and Corollary 9 are checked on every rendered frame; the rendered barrier filling attains δ^* within $1.2 \cdot 10^{-5}$ of the value range (the strict-extremum clearance); $F_T(0) = \tau^*$ holds exactly; and no hierarchy-preserving method falls below the width-free certificate.

7.2 RQ1: How common are hierarchy conflicts?

Under the automatic direction, $H > \theta$ in 34%, 24%, and 63% of the time steps of ERA5 1999/2000, ERA5 2013/14, and wildfire, with maxima of 25%, 16%, and 19% of the axis and medians of 12%, 7%, and 7% among conflict frames (Fig. 4). The space cost never exceeds 5%: some order could place the features almost exactly, and it is the hierarchy that prevents it. Conflicts are not an artefact of the threshold: for $\theta = 1\%$ and 5% the shares are 25–67% and 18–57%. Nor are they an artefact of the direction: across eight fixed directions the share ranges from 11% to 77%. And they are not an artefact of our width model: the width-free certificate exceeds θ in 31%, 22%, 63%, 50%, and 69% of the frames of the five datasets. τ_{free} is exact: a reference-sorted order is accepted only when it attains a proved lower bound; otherwise all orders are enumerated.

Prevalence depends on extraction. A conflict needs at least three leaves; frames average 3.4 leaves on ERA5 and 7.9 on fire. At a fixed 150 km ERA5 smoothing, chosen independently of enumeration limits, a certificate-only lower bound proves conflicts in at least 71% and 51% of the two winters (maximum 12 and 14 leaves). At 250/350 km, exact shares are 34/17% and 24/14%, respectively;




Figure 4: Per time step, space cost τ_{free} (grey) and hierarchy cost $H = \tau^* - \tau_{\text{free}}$ (orange) in % of the axis, on the three real datasets; dotted: θ . Positional distortion under the hierarchy is almost entirely hierarchy cost.

Method	ERA5	ERA5-14	Fire	Ring
TMTM	26.9 / 47.1	29.9 / 47.7	21.4 / 39.3	22.6 / 64.5
ST-MTM	21.2 / 41.3	23.4 / 53.3	21.4 / 35.7	7.3 / 34.7
A (tree kept)	19.2 / 23.1	15.9 / 29.9	16.1 / 32.1	24.2 / 31.5
R, $\kappa = 2\%$	17.3 / 24.0	18.7 / 29.9	12.5 / 42.9	20.2 / 22.6
R, $\kappa = 20\%$	13.5 / 18.3	13.1 / 23.4	6.2 / 31.2	10.5 / 12.9
R, $\kappa = \infty$	10.6 / 13.5	13.1 / 22.4	3.6 / 33.0	10.5 / 12.9
sorted, packed	10.6 / 14.4	13.1 / 22.4	3.6 / 33.0	8.9 / 12.1
exact projection	10.6 / 13.5	14.0 / 24.3	5.4 / 26.8	9.7 / 13.7
clear cases	104	107	112	124

Table 2: Geometric motion proxy, $k=2$: reversal / total error (reversal + miss) in % of clear 2-D pair-windows. A: CD-MTM with tree kept; R: relaxed. Largest δ^* in % of the value range for R at $\kappa = 2\% / 20\% / \infty$: ERA5 1.2 / 26.1 / 58.7; ERA5-14 1.2 / 19.5 / 28.6; fire 1.7 / 19.6 / 52.6; Ring 2.2 / 12.3 / 12.3.

for fire leaf caps of 8/10/12 they are 57/73/63%. Boundary minima also inflate prevalence: on the narrower window with boundary minima retained, the ERA5 shares are 58% and 60%. These are protocol-dependent rates, not general incidence estimates.

Simplification reduces conflicts by deleting features. Persistence simplification at 1% and 2% of the value range removes 28% and 34% of the ERA5 leaf instances (30% and 39% in 2013/14; 22% and 35% for fire), and conflicts remain in 24% and 19% of the ERA5 frames (13% and 7%; 57% and 33%). Only at 5% do conflicts fall to 3–10%, after about half of the pressure minima and 60% of the hot spots are gone. Flattening keeps every feature and pays in merge levels instead.

What existing maps could avoid. Under their best monotone calibration, TMTM and ST-MTM deviate from the reference by 11–29% of the axis on average on the real datasets, of which 2–6% is unavoidable for any merge-tree-preserving order; the remainder reflects that they optimize other goals than position (Fig. 1a).

7.3 RQ2: Geometric motion-proxy error

A reader of a static map may judge approach or separation by a pair’s vertical distance change. We simulate this geometric reading over $k=2$ steps (24 h for ERA5, four days for fire), with a 2% change threshold relative to each method’s occupied anchor range. We use two separate truths: change in 2-D extremum distance beyond 2% of the domain diagonal (Tab. 2), and change in projected-reference distance beyond 2% of the reference-axis extent. A *reversal* reads a clear change as its opposite, a *miss* as no change; their sum is the error. Each clear pair-window is one case. Paired differences use a circular moving-block bootstrap over time with $k+2$ -step blocks. If every anchor error is at most B_t at time t , then the displayed and projected pair-distance changes differ by at most $2(B_t + B_{t+k})$; this sufficient reading margin concerns the projection, not 2-D motion.

Existing maps and exact projection. Table 2 shows that existing maps reverse 21–30% of clear readings on real data, below but not far from the permutation null (32–39%). Even the exact-projection baseline, which draws every feature at its reference, reverses 11–14% on pressure and 5% on fire, because a single direction cannot show motion orthogonal to it. It is a reference point, not a bound: R at $\kappa = \infty$ on fire and ERA5 2013/14 and ST-MTM on Ring fall below it. Keeping the merge tree (A) lies 9 points above it on ERA5 1999/2000, consistent with the hierarchy cost of RQ1, and 2 points above it on 2013/14, where conflicts are rarer.

Projection versus layout loss. Against projected truth, R at $\kappa = 20\%$ reverses 5.8%, 0%, 2.8%, and 0% of clear cases on ERA5, ERA5-14, fire, and Ring, versus 13.5%, 13.1%, 6.2%, and 10.5% against 2-D truth. The projected clear-case counts are 120, 107, 107, and 140, respectively (the 2-D counts are in Tab. 2); denominators differ because the two truths define different clear events. Exact projection has zero projected error by construction. R still misses 5.8%, 4.7%, 15.0%, and 0.7% of projected cases because packing and the operating budget remain. The contrast shows why the reference task and 2-D task must be reported separately; their rates cannot be subtracted to estimate a projection effect.

Simple sorting baseline. The packed reference sort is competitive when topology is unrestricted: its 2-D reversal/error is 10.6/14.4%, 13.1/22.4%, 3.6/33.0%, and 8.9/12.1% (Tab. 2). Its largest attainable merge-level distortion is 58.7%, 28.6%, 52.6%, and 12.3% of each value range, compared with 26.1%, 19.5%, 19.6%, and 12.3% for R at $\kappa = 20\%$. On Ring it also has a lower reversal rate than R at the same largest cost. R’s contribution is therefore a disclosed, selectable position–topology trade-off, not dominance of simple sorting on the proxy.

Relaxation as a trade-off. Benefit and cost both depend on the merge-gap threshold κ . At $\kappa = 2\%$ the largest merge-level error stays below 2% of the value range, and reversals on pressure stay within 3 points of A. At $\kappa = 20\%$, reversals drop to 13–14% on pressure and 6% on fire, below ST-MTM by 7.7, 10.3, and 15.2 points (95% intervals $[-17.0, +2.0]$, $[-19.6, -1.0]$, $[-26.2, -3.9]$; the first is not significant, and fire has 8 blocks and is descriptive). The combined error drops by 23.1 and 29.9 points on pressure ($[-34.0, -12.7]$, $[-39.8, -20.0]$), because the relaxed map also misses fewer changes; on fire it drops by 4.5 points, which is not significant, as reversals turn into misses. The price is a largest merge-level error of 20–26% of the value range, with a mean over frames of 1.5–4%. At $\kappa = \infty$, R reaches the exact-projection baseline at a largest merge-level error of 29–59%. On Ring, where ST-MTM uses its published parameters, it reverses less (7.3% vs 10.5%, not significant) but misses far more (error 35% vs 13%).

Robustness. At $\kappa = 20\%$, using fixed x instead of the fitted direction gives reversal rates of 12–13% on the real datasets. Across three reader thresholds and four time windows, R reverses less often than ST-MTM in 16 of 18 settings, significantly in 9 (exploratory,

without multiplicity correction); a 5% threshold ties, and the 4-step ERA5-14 window is an exception. A grid covering the four published ST-MTM presets shows that settings with few reversals often freeze the layout (up to 95% misses); the best combined error remains above R's on all four datasets. Extraction matters: at 150 km smoothing ST-MTM's optimizer fails on ERA5 1999/2000, and with a coarser fire extraction (leaf cap 8) relaxation does not lower reversals. The proxy assumes exact feature identities and anchors; it does not measure human performance.

Does the certificate locate errors? A frame-level flag ($H > \theta$) marks most clear cases and does not predict where TMTM or ST-MTM reverse. A per-feature flag is more informative. We mark a pair when either feature deviates by more than θ under the best monotone calibration of its map (Proposition 6), so that the flag applies to any map. It catches 75–89% of the reversals of A on real data, and flagged cases are 1.8–3.2 times as likely to be reversed. It is not selective, marking 51–71% of clear cases, and unflagged cases still reverse in 6–13%. For R at $\kappa = 20\%$ few cases remain flagged (2–32%), and the flag catches 0–57% of the remaining reversals. The flag is a diagnostic, not a guarantee.

7.4 RQ3: What does fidelity cost, and is the relaxation optimal?

Figure 5 shows, per dataset, the exact frontier F , the certified bound Φ , and our layouts. Theorem 7 holds in all 656 rendered frames with broken contiguity, and no layout lies below Φ or F .

Cost. At $\kappa = 2\%$, relaxation lowers the mean reference error by 9–17% on pressure at 1.1–1.3 hPa of merge-level error with the optimal filling (1.6–2.2 hPa with the LCA filling), but by 45% on fire. At $\kappa = 20\%$ it lowers the error by 48–62% on pressure at 22–25 hPa. The cost of positional fidelity depends on the merge structure: fire conflicts stem from weak merges between neighbouring hot spots, pressure conflicts from strong saddles and shallow lows.

Optimality on relaxed frames. The selected order's $(\tau(\pi), \delta^*(\pi))$ lies within half a pixel of the exact frontier in 44/44, 99/102, and 106/110 relaxed frames at $\kappa = 2\%, 20\%, \infty$ (four datasets), never more than 0.7% of the axis above it. The threshold-and-prune policy's τ -optimal orders are on the frontier in every relaxed frame. Greedy one-node flattening stalls on plateaus, resolving only 88/110 conflicts at $\kappa = \infty$; flattening without pruning changes the tree in 73 rather than 44 frames at $\kappa = 2\%$. The aggregate points lie above F because κ thresholds merge gaps, whereas actual δ^* also includes leaf heights.

Tightness of the certified bound. $F - \Phi$ averages 0.3–1.2% of the axis over all breakpoints of Φ and is within half a pixel at 29–67% of them; it exceeds 10% in a few pressure and fire frames (at most 22%). Φ ignores the leaf-height term of Proposition 8, which can only widen this gap. The optimal filling halves d_{top} where only merge gaps matter (Ring, Gaussians) but saves little on pressure at large budgets, again because of shallow lows.

Hidden costs. Relaxation could harm properties that the reader proxy does not measure. Table 3 reports method-agnostic measures. On fire, tracked pairs swap vertical order in 3% of consecutive steps on R, against 8% and 33% on ST-MTM and TMTM; on ERA5, R (7%) lies between ST-MTM (5%) and TMTM (9%), and on Ring ST-MTM swaps fewer. Displacement-based stability is mixed: on pressure data R's anchors move more per step than ST-MTM's, as they follow the lows. R's 2-D distance preservation is worse than that of ST-MTM on pressure data (nearest-neighbour preservation 0.70 vs 0.88 on ERA5) and similar to that of exact projection (0.66), so much of this cost comes from reading a single direction; ST-MTM is the better choice when 2-D neighbourhoods matter more than positions along a direction. At $\kappa = 20\%$ the barrier filling changes 2–6% of the pixels, all in relaxed columns; unrelaxed columns are untouched.

Method	flips	jitter	stress	NN
TMTM	.092 / .334	.046 / .102	.317 / .458	.746 / .639
ST-MTM	.050 / .076	.027 / .053	.194 / .256	.877 / .753
A	.110 / .094	.101 / .054	.327 / .261	.731 / .658
R, $\kappa = 20\%$	.073 / .026	.089 / .022	.308 / .221	.700 / .630
R, $\kappa = \infty$	.083 / .010	.086 / .019	.296 / .219	.682 / .600
exact projection	.078 / .005	.086 / .017	.310 / .227	.660 / .589

Table 3: Hidden costs on ERA5 / fire: share of tracked pairs that swap vertical order between consecutive steps (flips), mean anchor displacement per step relative to the occupied range (jitter), Kruskal stress-1 and nearest-neighbour preservation (NN) of 1-D against 2-D extremum distances. Lower is better for the first three, higher for NN.

7.5 Case study: reading a winter with the certificate

Figure 6 shows how an analyst would use the certificate on ten days of the ERA5 winter. The hierarchy strip of the merge-tree-preserving map (a) exceeds θ in eight of the twenty steps, up to 25% of the axis, mostly between 1999-12-06 and 1999-12-11: on these days positional readings require caution. The witness triples explain why. At step 48 (1999-12-11 00 UTC), a deep low of approximately 985 hPa west of Scotland and one of 988 hPa over Scandinavia merge at 994.1 hPa, before either merges with a shallower 1007 hPa low over the Adriatic that lies between them along the reference direction (b). Any map that keeps the merge tree must place the 1007 hPa low outside the interval of the other two (c); this witness alone forces 7.8% of the axis, most of the certificate (8.4%). If the tracks matter, the relaxed map (d) removes the conflict, with an error of 0.8% of the axis at step 48. The topological price is explicit: even with the optimal filling, the west-Scotland/Scandinavia merge reads at 1006.7 rather than 994.1 hPa, an error of 12.6 hPa. An analyst studying the depth of that shared trough should keep the tree-preserving view.

7.6 RQ4: Does the certificate transfer?

Dense pixel displays often order rows by a dendrogram. We ordered 1024 active 10-km fire cells, one pixel row each, by an agglomerative clustering of their fire time series, the hierarchy behind hierarchical orderings of pixel displays [16]. The reference is each cell's position along the principal geographic axis, scaled linearly to the display height. The certificate shows that any order consistent with this dendrogram must displace some cell by at least 67% of the display height, whereas the geographic order itself stays within 21%. With a rank-based reference the bound is 78–82%. The hierarchy cost is 47–48% for average, complete, and Ward linkage, versus 1.4% for a spatial Ward dendrogram and 32.5% for a spatial single-linkage dendrogram. With optimal leaf ordering, cells sit up to 88% (median 14%) away from geography, and 8% of the rows exceed the certificate (supplementary dendrogram figure). The certificate thus answers a design question before anyone reads geography off a clustered display: here, “row position approximates location” is untenable, and a linked map view or a spatially constrained hierarchy is needed. The computation takes 0.1 s.

7.7 Computational scope and reproducibility

On recorded low-degree trees, the certificate takes about 2 ms per small frame at 1024 pixels and 3.6 s for 4096 leaves at 16384 pixels. In a fixed 12-leaf, 96-pixel stress test, flattening the root from degree 2 to 12 raises subset states from 3 to 4095 and DP time from 0.001 to 0.135 s. Frontiers enumerate $n!$ orders; layout selection also enumerates orders and solves QPs. Large-scale evidence therefore applies to diagnosis, not the complete relaxation pipeline. The supplement includes source, independent optimization checks, experiment entry points, selected JSON results, and a hash manifest; public raw inputs are documented separately.




Figure 5: Exact frontier F (blue), certified bound Φ (dashed), and R at four κ values (filled: optimal filling; open: LCA filling). x : largest topological error over frames; y : mean fixed-order optimum $\tau(\pi)$, excluding layout slack β .




Figure 6: ERA5 winter case. (a) Tree-preserving map: blue bars and orange rings mark witness triples in conflict steps. (b) Step 48 field and projected extrema. (c) References versus tree-preserving anchors: the orange low cannot lie between the blue lows. (d) Relaxed map ($\kappa = 20\%$) with position-error and topological-cost strips.

8 Discussion and Limitations

Choosing a view. For comparing trough depth or merge structure (T1), the tree-preserving view retains the relevant scalar relations. For positions along the reference (T2), the certificate shows when exact placement is impossible and witnesses explain why. For projected motion comparisons (T3), relaxed views can reduce geometric proxy error, but the topology strip must be consulted: the winter case trades a positional conflict for a 12.6 hPa merge error. When 2-D neighbourhoods matter, ST-MTM can preserve them better than CD-MTM (Tab. 3); a linked spatial view remains useful. The result is a selectable trade-off with measured costs, rather than a universal preferred layout.

Limits of the theory. One reference coordinate cannot encode orthogonal motion; radial or path references need separate evaluation. The optimal filling and frontier assume true extrema values. The certificate DP is pseudo-polynomial in resolution and exponential in node degree; layout selection and the exact frontier enumerate orders. Complexity for unbounded degree remains open. Φ ignores leaf heights and becomes expensive after flattening. Relaxation has no approximation guarantee, and κ thresholds merge gaps rather than actual topological cost.

Limits of the evaluation. The geometric proxy does not establish human performance. The direction is fitted on the evaluated tracks, though the fixed-direction control retains the reversal gain. ST-MTM was reimplemented from its description; we

could not compare with its original code. ERA5 comes from a public mirror and is not byte-identical to the original files; window and boundary-minimum choices affect conflict rates (Sec. 7.2). At 150 km smoothing we report certificate lower bounds only, not full relaxation or motion results. The fire background term was not varied.

Next validation steps. A task-based study should test whether the disclosures help readers choose between views, measuring merge-level accuracy and projected motion judgments separately. Additional extraction protocols and directions fitted on held-out data would test generalization.

9 Conclusion

CD-MTM combines positional certificates, optimal value-preserving filling, and visual disclosure to make the position-topology trade-off explicit. The certificate identifies unavoidable displacement; the frontier measures what improved placement costs in merge levels. In the selected protocols, 24–63% of frames conflict, and relaxation reduces several geometric motion-proxy errors at substantial topological cost. These results support diagnosis and informed layout selection; human benefits and scalable full-layout optimization remain open.

Data and figure credits. Figures: authors. Fire data: Franke et al. [16] (CC BY 4.0); ERA5: Copernicus [18]; coastlines: Natural Earth (public domain).

References

- [1] M. Aupetit. Visualizing distortions and recovering topology in continuous projection techniques. *Neurocomputing*, 70, 2007. 1, 2
- [2] Z. Bar-Joseph, D. K. Gifford, and T. S. Jaakkola. Fast optimal leaf ordering for hierarchical clustering. *Bioinformatics*, 17(Suppl 1):S22–S29, 2001. 2
- [3] K. Beketayev, G. H. Weber, D. Morozov, A. Abzhaynov, and B. Hamann. Geometry-preserving topological landscapes. In *Proc. WASA (SIGGRAPH Asia Workshop)*, pp. 155–160, 2012. 2
- [4] T. Beurskens, S. van den Broek, A. Simons, W. Sonke, K. Verbeek, T. Ophelders, M. Hoffmann, and B. Speckmann. ParkView: Visualizing monotone interleavings. In *Proc. IEEE Pacific Visualization Conference*, pp. 118–127, 2025. doi: 10.1109/PacificVis64226.2025.00018 2
- [5] J. Buchmüller, D. Jäckle, E. Cakmak, U. Brandes, and D. A. Keim. MotionRugs: Visualizing collective trends in space and time. *IEEE Transactions on Visualization and Computer Graphics*, 25(1), 2019. 1, 2
- [6] L. Bulteau, P. Gambette, and O. Seminc. Reordering a tree according to an order on its leaves. In *Proc. CPM*, 2022. 2
- [7] H. Carr, J. Snoeyink, and M. van de Panne. Simplifying flexible iso-surfaces using local geometric measures. In *Proc. IEEE Visualization*, pp. 497–504, 2004. doi: 10.1109/VISUAL.2004.96 2
- [8] V. Chepoi and B. Fichet. ℓ_∞ -approximation via subdominants. *Journal of Mathematical Psychology*, 44(4):600–616, 2000. doi: 10.1006/jmps.1999.1270 4
- [9] M. Diaz, J. Heredia, N. Ferreira, and J. Poco. Spatiotemporal merge tree maps: A topology and geometry aware visualization for temporal scalar data, 2026. Preprint, SSRN 6604235. doi: 10.2139/ssrn.6604235 1, 2, 6
- [10] H. Doraiswamy, J. Tierny, P. J. S. Silva, L. G. Nonato, and C. Silva. TopoMap: A 0-dimensional homology preserving projection of high-dimensional data. *IEEE Transactions on Visualization and Computer Graphics*, 27(2):561–571, 2021. doi: 10.1109/TVCG.2020.3030441 2
- [11] T. Dwyer, Y. Koren, and K. Marriott. IPSep-CoLa: An incremental procedure for separation constraint layout of graphs. *IEEE Transactions on Visualization and Computer Graphics*, 12(5):821–828, 2006. 2
- [12] H. Edelsbrunner, D. Letscher, and A. Zomorodian. Topological persistence and simplification. *Discrete & Computational Geometry*, 28(4):511–533, 2002. doi: 10.1007/s00454-002-2885-2 2
- [13] M. B. Eisen, P. T. Spellman, P. O. Brown, and D. Botstein. Cluster analysis and display of genome-wide expression patterns. *Proceedings of the National Academy of Sciences*, 95(25):14863–14868, 1998. 1, 2
- [14] J. Eldridge, M. Belkin, and Y. Wang. Beyond Hartigan consistency: Merge distortion metric for hierarchical clustering. In *Proc. Conference on Learning Theory (COLT)*, vol. 40 of PMLR, pp. 588–606, 2015. 2
- [15] M. Farach, S. Kannan, and T. Warnow. A robust model for finding optimal evolutionary trees. *Algorithmica*, 13(1–2):155–179, 1995. doi: 10.1007/BF01188585 4
- [16] M. Franke, H. Martin, S. Koch, and K. Kurzhals. Visual analysis of spatio-temporal phenomena with 1D projections. *Computer Graphics Forum*, 40(3):335–347, 2021. 2, 6, 8, 9
- [17] E. Gasparovic, E. Munch, S. Oudot, K. Turner, B. Wang, and Y. Wang. Intrinsic interleaving distance for merge trees. *La Matematica*, 2024. 2
- [18] H. Hersbach et al. The ERA5 global reanalysis. *Quarterly Journal of the Royal Meteorological Society*, 146, 2020. 5, 9
- [19] E. Hovmöller. The trough-and-ridge diagram. *Tellus*, 1(2), 1949. 1
- [20] D. H. Huson. Displacement-optimized tanglegrams for trees and networks. *Molecular Biology and Evolution*, 43(3), 2026. 2
- [21] D. A. Keim. Designing pixel-oriented visualization techniques: Theory and applications. *IEEE Transactions on Visualization and Computer Graphics*, 6(1):59–78, 2000. 2
- [22] G. Kindlmann and C. Scheidegger. An algebraic process for visualization design. *IEEE Transactions on Visualization and Computer Graphics*, 20(12), 2014. 2
- [23] J. Klawitter, F. Klesen, J. Y. Scholl, T. C. van Dijk, and A. Zaft. Visualizing geophylogenies – internal and external labeling with phylogenetic tree constraints. *Journal of Graph Algorithms and Applications*, 29(1), 2025. 1, 2
- [24] W. Köpp and T. Weinkauff. Temporal treemaps: Static visualization of evolving trees. *IEEE Transactions on Visualization and Computer Graphics*, 25(1), 2019. 2
- [25] W. Köpp and T. Weinkauff. Temporal merge tree maps: A topology-based static visualization for temporal scalar data. *IEEE Transactions on Visualization and Computer Graphics*, 29(1):1157–1167, 2023. 1, 2, 6
- [26] S. Li and H. Wang. Separating overlapped intervals on a line. *Journal of Computational Geometry*, 10(1):281–321, 2019. 2
- [27] S. Liu, Y. Wu, E. Wei, M. Liu, and Y. Liu. StoryFlow: Tracking the evolution of stories. *IEEE Transactions on Visualization and Computer Graphics*, 19(12):2436–2445, 2013. doi: 10.1109/TVCG.2013.196 2
- [28] J. Lukaszczuk, G. Weber, R. Maciejewski, C. Garth, and H. Leitte. Nested tracking graphs. *Computer Graphics Forum*, 36(3), 2017. 2
- [29] A. McNutt, G. Kindlmann, and M. Correll. Surfacing visualization mirages. In *Proc. CHI*, 2020. 1, 2
- [30] W. Meulemans. An algorithmic perspective on information visualization, 2026. arXiv:2607.29360. 2
- [31] D. Morozov, K. Beketayev, and G. H. Weber. Interleaving distance between merge trees. In *Proc. TopoInVis*, 2013. 2
- [32] E. Munch and A. Stefanou. The ℓ_∞ -cophenetic metric for phylogenetic trees as an interleaving distance. In *Research in Data Science*. Springer, 2019. 2, 4
- [33] Q. H. Nguyen, P. Eades, and S.-H. Hong. On the faithfulness of graph visualizations. In *Proc. IEEE Pacific Visualization Symposium*, pp. 209–216, 2013. doi: 10.1109/PacificVis.2013.6596147 2
- [34] L. G. Nonato and M. Aupetit. Multidimensional projection for visual analytics: Linking techniques with distortions, tasks, and layout enrichment. *IEEE Transactions on Visualization and Computer Graphics*, 25(8):2650–2673, 2019. doi: 10.1109/TVCG.2018.2846735 2
- [35] S. Nusrat and S. Kobourov. The state of the art in cartograms. *Computer Graphics Forum*, 35(3), 2016. 2
- [36] M. Pont, J. Vidal, J. Delon, and J. Tierny. Wasserstein distances, geodesics and barycenters of merge trees. *IEEE Transactions on Visualization and Computer Graphics*, 28(1), 2022. 2
- [37] J. Rauscher, F. L. Dennig, U. Schlegel, D. A. Keim, and T. Schreck. Visual boosting techniques for spatiotemporal dense pixel visualizations. In *Proc. EuroVA*, 2026. 2
- [38] J. Rauscher et al. Visually assessing 1-D orderings of contiguous spatial polygons. *Computer Graphics Forum*, 44, 2025. 2
- [39] B. Rieck and H. Leitte. Persistent homology for the evaluation of dimensionality reduction schemes. *Computer Graphics Forum*, 34(3):431–440, 2015. doi: 10.1111/cgf.12655 2
- [40] M. C. A. Stolk, J. J. H. M. Wulms, and K. A. B. Verbeek. GroupRugs: Visual summaries for groups in collective movement data. In *Proc. PacificVis*, pp. 128–137, 2025. 2
- [41] J. Tierny and V. Pascucci. Generalized topological simplification of scalar fields on surfaces. *IEEE Transactions on Visualization and Computer Graphics*, 18(12):2005–2013, 2012. doi: 10.1109/TVCG.2012.228 2
- [42] G. Valdrighi, N. Ferreira, and J. Poco. MoReVis: A visual summary for spatiotemporal moving regions. *IEEE Transactions on Visualization and Computer Graphics*, 30(4), 2024. 2
- [43] N. van Beusekom, W. Meulemans, B. Speckmann, and J. Wood. Data-spatial layouts for grid maps. In *Proc. GIScience*, vol. 277 of LIPIcs, 2023. 2
- [44] J. Venna and S. Kaski. Neighborhood preservation in nonlinear projection methods: An experimental study. In *Proc. ICANN*, 2001. 1, 2
- [45] M. Wattenberg, F. Viégas, and I. Johnson. How to use t-SNE effectively. *Distill*, 2016. 1, 2
- [46] G. H. Weber, P.-T. Bremer, and V. Pascucci. Topological landscapes: A terrain metaphor for scientific data. *IEEE Transactions on Visualization and Computer Graphics*, 13(6), 2007. 2

- [47] J. Wood and J. Dykes. Spatially ordered treemaps. *IEEE Transactions on Visualization and Computer Graphics*, 14(6):1348–1355, 2008. [2](#)
- [48] J. Wulms, J. Buchmüller, W. Meulemans, K. Verbeek, and B. Speckmann. Stable visual summaries for trajectory collections. In *Proc. PacificVis*, 2021. [2](#)
- [49] L. Yan, T. B. Masood, F. Rasheed, I. Hotz, and B. Wang. Geometry-aware merge tree comparisons for time-varying data with interleaving distances. *IEEE Transactions on Visualization and Computer Graphics*, 2022. [2](#)
- [50] L. Zhou, C. R. Johnson, and D. Weiskopf. Data-driven space-filling curves. *IEEE Transactions on Visualization and Computer Graphics*, 27(2):1591–1600, 2021. [2](#)